



Lectures: Data Structures

Programming Systems and Tools

Vectors

State University of New York at Binghamton

May 2, 2024

BINGHAMTON | THOMAS J. WATSON COLLEGE OF
UNIVERSITY | ENGINEERING AND APPLIED SCIENCE

3 / 121

Motivation



Suppose we have a mutable data set with the following requirements:

- Varies in size from 10,000 to 10,000,000 elements
- Requires random access
- Can we use a typical C array?

1 Vectors

2 Lists

3 Stacks and Queues

4 Trees

5 Heaps

6 Priority Queues

4 / 121

2 / 121

Definition



Vectors: Dynamic Arrays



We will implement struct `Vector` using this definition:

```
1 typedef uint64_t Data;
2 typedef struct Vector {
3     size_t size; // index of last element in the array
4     size_t capacity; // the total number of elements that can be stored
5     Data *array; // pointer to the storage space
6 } Vector;
7
8 Vector *newVector(size_t size);
9 void deleteVector(Vector *v);
10 int printVector(const void *v);
11
12 Data *at(Vector *v, size_t index);
13 size_t capacity();
14 void clear(Vector *v);
15 Data *find(Vector *v, Data d);
16 Data *insert(Vector *v, Data d, size_t idx);
17 Vector *resize(Vector *v, size_t count);
18 Data *remove(Vector *v, size_t index);
```

7 / 121

Initializing Vectors

```
Vector *newVector(size_t size);
```

- 1 `v ← (Vector *) malloc(sizeof(Vector))`
- 2 `v.capacity ← size`
- 3 `v.size ← 0`
- 4 `v.array ← (Data *) malloc(sizeof(Data) * v.capacity)`
- 5 Use `memset` to zero-initialize the array storage
- 6 `return v`

5 / 121

Vectors: Dynamic Arrays



Pros

- Allocates memory only as needed

- Allows dynamically-sized arrays

- Resizing penalty can be amortized

- Use `memset` to zero-initialize the array storage

Cons

- Memory and performance intensive

- Still wastes space, particularly at larger array sizes

- Members must be a single type (not a "bag")

8 / 121

6 / 121



What happens if the size argument is less than v.size?

- Allow the resize, but some data is lost
- Disallow the resize, return a pointer to the original Vector
- Allow realloc() default behavior:
 - The contents will be unchanged in the range from the start of the region up to the minimum of the old and new sizes
 - From the minimum size (new capacity) to the maximum size (old capacity), the data is not copied

Resizing the Vector



Resizing the Vector

```
Vector *resize(Vector *v, size_t count);
```

We will eventually need to grow our dynamic array's storage space. Two strategies:

- Every time we add, we *incrementally* expand the array (e.g. if size was 10, we resize to 11, 12, 13, etc.)
 - Keeps array exactly large enough for current needs
 - Growth operations require heap reallocation on every insertion ... **no amortization!**
- If length equals the size, we *geometrically* increase the size to the next power of two (e.g. if size was 10, we resize to 16, 32, 64, etc.)
 - Array will outsize the current needs, particularly at large array sizes.
 - Growth operations can be amortized



```
Data *insert(Vector *v, Data d, size_t idx);
```

Code Walkthrough

```
1 Data *insert(Vector *v, Data d, size_t idx);
2
3     if (idx > v.len) {
4         v.array[idx] = d;
5     }
6     else {
7         if (v.array[idx] == 0) {
8             v.array[idx] = d;
9         }
10    }
11 }
```

Reallocation in the C Standard Library



Reallocation in the C Standard Library



```
1 #include <stdio.h> // printf
2 #include <stdlib.h> // malloc, realloc, free
3 #include <string.h> // strcpy, strcat
4 int main () {
5     char *str;
6
7     /* Initial memory allocation */
8     str = (char *) malloc(6);
9     strcpy(str, "Hello");
10    printf("String = %s, Address = %p\n", str, str);
11
12    /* Reallocating memory */
13    str = (char *) realloc(str, 14);
14    strcat(str, " World!");
15    printf("String = %s, Address = %p\n", str, str);
16
17    /* Undefined freeing */
18    str = (char *) realloc(str, 0);
19    //free(str);
20    return(0);
21 }
```

15 / 121

13 / 121

Reading from a Vector



```
Data *at(Vector *v, size_t index);
```

Reading from vectors is just like using the indexing operator on arrays, except:

- Instead of returning the actual value, we return a pointer; this allows us to have a failure state...
- We must check bounds: if the index is larger than the bounds of the vector, return NULL
 - *No!* ... We're not causing a buffer overread (the memory's allocated, right?)
 - *Yes!* ... We're violating the intent of the data structure
 - **We'll consider it a problem and return a nullptr.**

```
1 #include <stdio.h>
2 void *realloc( void *ptr, size_t new_size );
```

realloc() reallocates the given area of memory.

- It must be previously allocated by malloc(), calloc() or realloc() and not yet freed with a call to free() or realloc().
- Reallocation happens by one of two strategies:
 - Expanding or contracting the existing area pointed to by ptr, if possible. The contents of the area remain unchanged up to the lesser of the new and old sizes. If the area is expanded, the contents of the new part of the array are undefined.
 - Allocating a new memory block of size new_size bytes, copying memory area with size equal the lesser of the new and the old sizes, and freeing the old block.
 - If there is not enough memory, the old memory block is not freed and null pointer is returned.

Experiment: Using realloc()



1 Create a file **realloc.c**

- 2 Allocate a 6 byte-wide character buffer on the heap using malloc().
- 3 Use strcpy() to copy the string "Hello" into the buffer
- 4 Print the address of the buffer, and the string in the buffer.
- 5 Use realloc() to change the size of the original buffer to 14 bytes.
- 6 Use strcat() to append the string " World!" to the original buffer.
- 7 Again, print the address of the buffer and the string in the buffer.
- 8 Call realloc() a second time, this time with size 0. Do not call free.
- 9 Compile the program and run it in Valgrind. Is there a memory leak?

16 / 121

14 / 121

Clearing a Vector



Removing from a Vector



```
void clear(Vector *v);
```

Clearing the vector is just iterating across the entire vector, and removing the value at each index:

1 For index i in range 0 ... v.size - 1: v.array[i] \leftarrow 0

2 v.size = 0

Alternative option: memset from index 0 to size ... Much faster!

■ Should we resize the array also? We do not. Why?

- Relies on memory allocator not causing a failure
- Expensive to call realloc()
- Don't quietly change state on the user.

```
Data *remove(Vector *v, size_t index);
```

- 1** Actual deletion by deallocation of that part of vector memory
 - Expensive and must reallocate space carefully
 - What if item deleted is not last item?
- 2** Set v.array[index] to an "empty" value
 - Most efficient option
 - Doesn't cause problems with memory management or external references
 - Need to reserve an "empty" value ... no "None" value in C
- 3** Remove the item and move items at higher indices up to size-1.
 - Most accurately models behavior of bounds checked buffers
 - Could cause expensive iteration to move items leftward!
 - Dangerous to mutate storage while also iterating across it!

19 / 121

17 / 121

Removing from a Vector



What if another variable outside the struct has a reference for an index into the array?

- Vector cannot stop unauthorized access into the array
- Unauthorized access might not cause .size or .capacity to be properly updated

Lists

20 / 121

18 / 121

Vectors Hide Underlying Complexity

Motivation



- Suppose we have a dynamic vector `v` with
`.array={1, 4, 10, 19, 6}`, `.capacity=5`.
- Then suppose we want to call `insert(v, 7, 2)`. What vector operations have to happen?
 - 1 `resize(v, 8)`:
 - 2 Traverse heap free list to find a new array storage: $O(n)$
 - 3 Acquire new heap storage: $O(1)$
 - 4 Copy array contents from old heap storage to new: $O(n)$
 - 5 Deallocate old heap storage: $O(1)$
 - 6 Update `v.capacity = 8`: $O(1)$
 - 7 Copy 3 array items to next index: $O(n)$
 - 8 Insert 7 at index 2: $O(1)$
 - 9 Update `v.size = 6`: $O(1)$

21 / 121

Linked Lists Remove Underlying Complexity

Vectors Hide Underlying Complexity



- Vectors have problems C dynamic vectors are:
 - **Inflexible:** Vectors store in statically sized arrays, meaning once the size is allocated the original array cannot change.
 - **Wasteful:** Over-allocation happens “in case you need it”
 - **Cumbersome:** Insertion requires resizing, resizing requires reallocation.

21 / 121

- Suppose we have a dynamic vector `v` with
`.array={1, 4, 10, 19, 6}`, `.capacity=5`.
- Then suppose we want to call `insert(v, 7, 2)`. What vector operations have to happen?
 - The problem with arrays is that they must be contiguous in memory and statically sized.
 - Resizing is entirely dependent on the current state of memory.
 - What if we could link non-contiguous segments of memory?
 - Insertion and deletion are less cumbersome – we just link to the subsequent or previous segment.
 - Resizing, inserting or deleting are a matter of changing links between segments.

23 / 121

22 / 121

What Goes in the Node?

Allowing mixed types in a List

Node Objects

- We will create our list using objects called **nodes**.

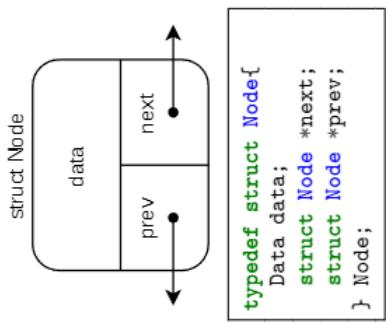
- Each node represents one piece of linked data and is comprised of two or three members:

- Its data payload, and either:

- A pointer to the next node only (*singly linked*)
 - A pointer to both the previous and next nodes (*doubly linked*)
- Every node contains the address of the next node and previous node with two exceptions:

- The **head** contains null in its previous pointer.
- The **tail** contains null in its next pointer.

```
1 #include <stdint.h>
2 /* typedef option */
3 #typedef uint64_t Data;
4
5 /* payload struct option */
6 typedef struct Data {
7     union {
8         uint64_t uint64;
9         int64_t int64;
10        } data;
11    enum type { UINT64, INT64 };
12    } Data;
```



What Goes in the Node?

What Goes in the Node?



26 / 121

What Goes in the Node?

Allowing mixed types in a List



24 / 121

Pros of struct Data

- We can have multiple payload types throughout the program
- We can have mixed payload types, even within the same data structure!

Cons of struct Data

- Harder to implement
- Must access the payload through the union
- Uses more storage space
- Must keep Data's enumeration in sync

Pros of typedefs

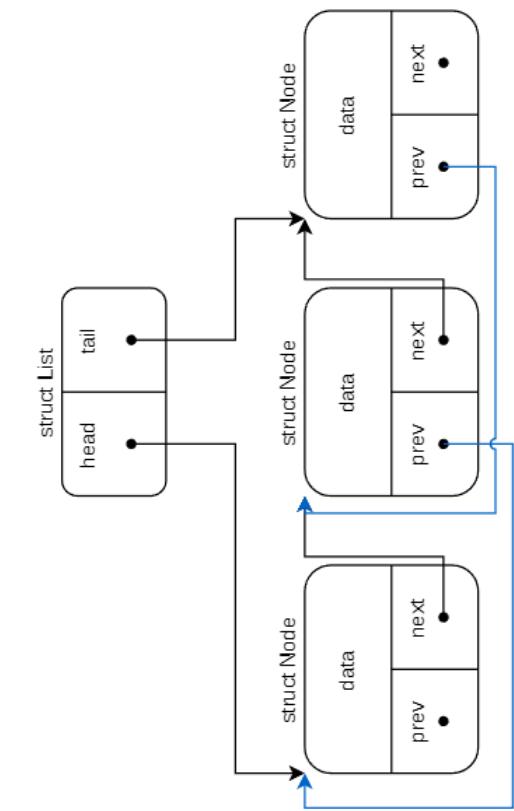
- Every Data member has to be same type throughout the program
- Must keep uses of Data synchronized to type

- No matter what kind of data is stored in the Data member, the algorithms for our List's "CRUD" (Create, Read, Update, Delete) operations will be the same.
- In other languages, we can use *templates* to hold the data, making the data structure type agnostic.
- In C, we *could* make our Nodes able to hold multiple types of data by typedefs or payload structs

27 / 121

Doubly Linked List

List Objects



- A linked list object's only *required* member is a pointer to the head of the list
- For a doubly-linked list, we can add a pointer to the tail for convenience, and a count of the number of linked Nodes in the List
- For an empty list, we set the head (and tail) pointers to `NULL`, and the size to 0
- **Do not use the head or tail pointers to traverse the list.** Copy them into another pointer first. ... **Why?**
- We don't include any data **payload** in the list object itself ... **Why?**

30 / 121

28 / 121

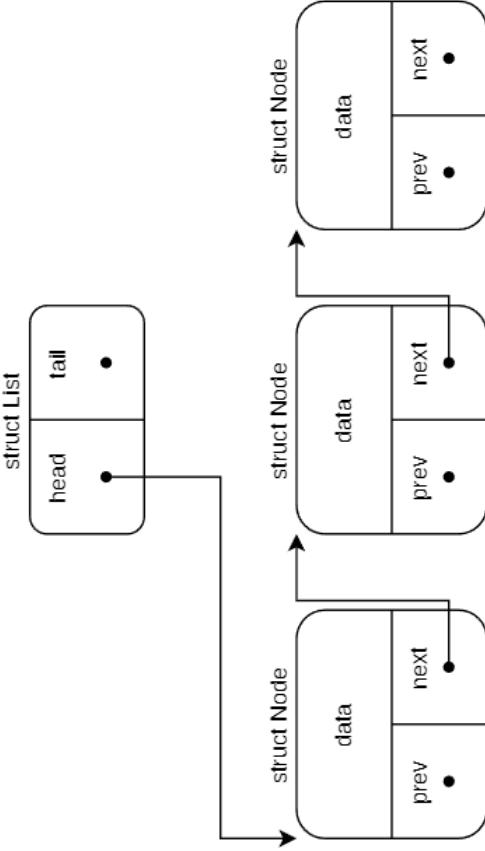
List Object Definition



We will implement the List with the following definition:

```
1 #ifndef DATA
2 #define DATA
3 #typedef uint64_t Data;
4 #endif // DATA
5
6 #ifndef LIST_H
7 #define LIST_H
8
9 #typedef struct Node {
10     Data data;
11     struct Node *next, *prev;
12     Node;
13 }
14 #typedef struct List {
15     Node *head, *tail;
16     size_t size;
17 } List;
18
```

Singly Linked List



31 / 121

29 / 121

Copying a Linked List



List Object Definition



```
1 Node *newNode(Data d);
2 void deleteNode(Node *n);
3 int printNode(const void *n);
4
5 List *newList();
6 void deleteList(List *l);
7 int printList(const void *l);
8
9 List *append(List *l, Node *n);
10 void clear(List *l);
11 Node *find(List *l, Data d);
12 List *insert(List *l, Node *n, Node *pos);
13 size_t length(List *l);
14 List *remove(List *l, Node *pos);
15 void traverse(List *l, int (*func)(const void *));
```

What happens when we use the assignment operator with two pointers to List objects? **struct List** new_list = old_list;
Shallow Copy Only copy the pointer address. Both list pointers point to the same list

Deep Copy Copy the sub-objects from the old object into the new object, including their values. Each list pointer points to one of two identical but separate lists
C will do a shallow copy with the assignment operator. We must implement a **List *copyList(List *l)** function to get deep copy through list traversal ... This is an exercise to the reader.

- All of the object management functions work as with Vectors:
- The new function returns a pointer to the newly created object
 - The delete function frees the memory of the object pointed to by the argument. In the case of the **deleteList()** function, it must also delete all the Nodes in the List first.

32 / 121

Getting the length of a List



Copying a Linked List



- What happens when we use the assignment operator with two pointers to List objects? **struct List** new_list = old_list;
Shallow Copy Only copy the pointer address. Both list pointers point to the same list
Deep Copy Copy the sub-objects from the old object into the new object, including their values. Each list pointer points to one of two identical but separate lists

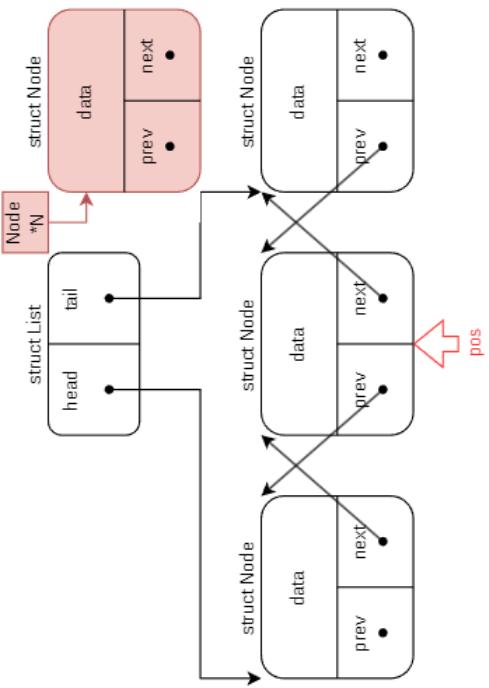
Two options:

- Traverse the List, counting the number of nodes we encounter until we reach either a NULL pointer or the tail node: $O(n)$
- Update the List.size field after each **append()**, **clear()**, **insert()**, **remove()** operation: $O(1)$ amortized

Adding Nodes to the List



Getting a Node from the List



```
Node *find(List *l, Data d);
```

- 1 Create an iterator pointer and initialize it to `l->head`
- 2 Check the contents of the node for a match,

- if matched, return
- If no match, advance the iterator to `iter->next`

- 3 If `iter == NULL`, break and return

Return values:

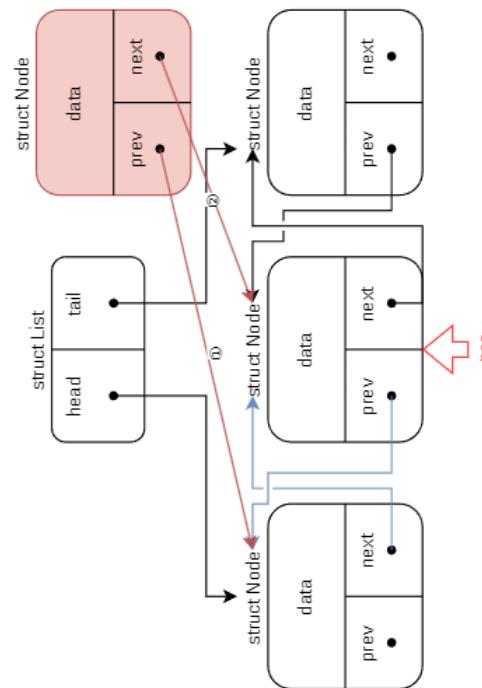
- We found the first node with the desired contents: return a pointer to that node
- We go past the tail of the list to a `nullptr`: we return a `NULL` pointer to indicate we couldn't find the right node
 - Either way, `iter` points to the value we want to return!

35 / 121

Adding Nodes to the List



Adding Nodes to the List



```
List *insert(List *l, Node *n, Node *pos);  
List *append(List *l, Node *n);
```

- The difference between the two functions is that `*insert_node` inserts `n` before `pos`. `List *append_Node(List *l, Node *n)` inserts the node after the last node in the List (i.e. at a new `l->tail`).

- When inserting, connect the new node to the list before updating pointers already in the lists' nodes.
- Once the new node is properly connected, it can be used to update the existing List node pointers.

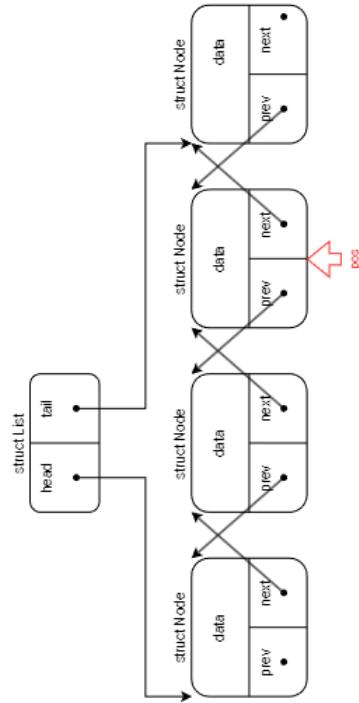
- Failure to properly connect the new node will cause a memory leak due to lost nodes, and/or segmentation faults when traversing the list.

37 / 121

36 / 121

Removing Nodes from the List

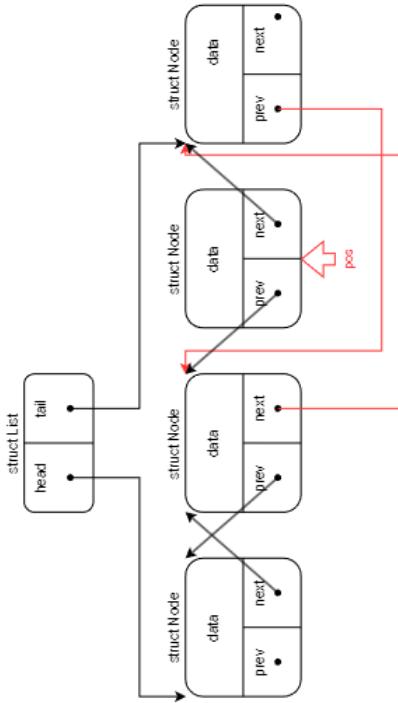
Adding Nodes to the List



39 / 121

Removing Nodes from the List

Removing Nodes from the List

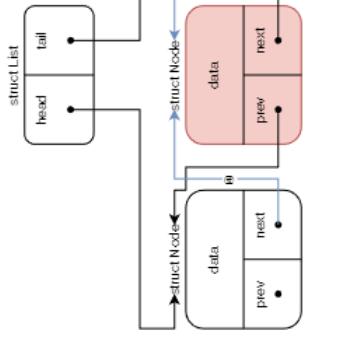


37 / 121

```
List *remove(List *l, Node *pos);
```

- 1 Find the node before pos if singly-linked (the predecessor, pred). The node after (the successor, succ) can be found using pos itself.
- 2 Update pred->next to point to successor, succ->prev to point to predecessor.
- 3 Delete the node pointed to by pos to prevent a memory leak and return a pointer to the newly updated list.

39 / 121

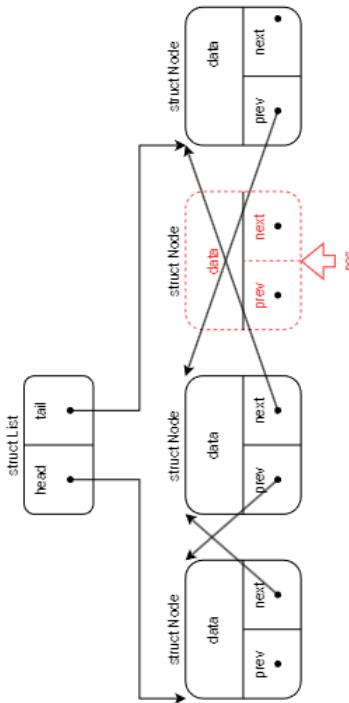


38 / 121

Abstract Data Types

B Removing Nodes from the List

- Data Structures (DS) are the implementation that manage data
- Abstract Data Types (ADT) describe behavior, not implementation
 - A DS is an implementation of an ADT
 - DS should have a standardized interface, but the underlying implementation can vary greatly



41 / 121

Stacks

Stacks are a sequential collection:

Sequential Order of element insertion matters

Collection Contains items of the same type, forming a grouping

- The purpose of a Stack ADT is to describe how the data is accessed, not how the data is stored!
- Two popular Stack DS implementations are linked list and array
 - The Stack DS implementations provide the ADT's sequential access control through encapsulation

39 / 121

B

Stacks and Queues

42 / 121

Examples of Stacks

Stack Access Control



What are some examples of stacks in computer science?

- Function call stack
- Exception handling/unwinding
- Undo/Redo commands
- Syntax checking - are brackets properly matched?
- Prefix notation calculators ("RPN")



- Stacks are **Last In, First Out (LIFO)**
- Items are removed in the reverse order of their insertion.
- Think of stacks as being like a pile of plates: to get to a lower plate, we have to lift the plates above it.
- Stacks limit the user to accessing only the newest element.

44 / 121

Stack Update Operations

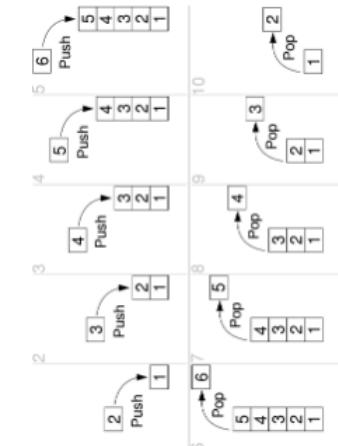
Examples of Stacks



What are some examples of stacks in computer science?

- Push
 - Inserting an item onto the stack
 - Imagine pushing down on everything already on the stack, one place farther from the top
- Pop
 - Removing from the stack
 - Each element remaining on the stack moves one place closer to the top
 - To access an element below, you must pop every element above it
 - `top()` or `peek()` get the top item like `pop()` but without removing it from the stack

43 / 121



45 / 121

44 / 121

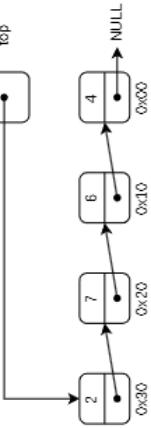
Stack Implementations

Stack Implementations



Pros of Array Based

- Easier memory management



Cons of Array Based

- Must have a static, maximum size or a dynamic vector storage

46 / 121

Stack Interface



Stack Implementations



```
1 #ifndef DATA_T
2 #define DATA_T
3 #typedef uint64_t Data;
4 #endif // DATA_T
5
6 #ifndef STACK_H
7 #define STACK_H
8
9 typedef struct Node {
10     Data value;
11     struct Node *next;
12 } Node;
13
14 typedef struct Stack {
15     Node *top;
16     size_t size;
17 } Stack;
```

Linked List Based

- 'top' is a pointer to the head of the list
- Stack grows at the head of the list

Pros of Linked List Based

- Internal storage is dynamically allocated → always right-sized

Cons of Linked List Based

- Push and pop have increased memory management

Queue Access Control

Stack Interface



- Queues are **First In, First Out** (FIFO)

- Items are removed in the order of their insertion.
- Think of queues as being like a line at the airport: each element must wait until all other elements that arrived first are processed.
- Queues limit the user to accessing only the oldest element.



```
1 // Not shown: new, delete, print methods
2 void clear(Stack *s);
3 size_t size(Stack *l);
4 Node *top(Stack *s);
5 Node *pop(Stack *s);
6 Node *push(Stack *s, Data d);
7 void traverse(Stack *s, int (*func)(const void *));
```

50 / 121

50 / 121

Examples of Queues



Queues

What are some examples of queues in computer science?

Sequential Order of element insertion matters

Collection Contains items of the same type, forming a grouping

- Like stacks, queues provide restricted access control through encapsulation
- Unlike stacks, queues are First In First Out (FIFO)

53 / 121

51 / 121

Queue Implementations



Examples of Queues



Array Based

- 'head' and 'tail' are pointers to the front and end of the queue in the array
- Head of the queue is at a lower address than the tail of the queue

Pros of Array Based

- Easier memory management

Cons of Array Based

- Requires array to have a static, maximum size or a dynamic vector

Queue Implementations



55 / 121

- What are some examples of queues in computer science?
 - I/O buffers
 - Operating system pipes
 - Multithreading synchronization primitives
 - First Come, First Serve (FCFS) process scheduling
 - Cryptographic block and stream ciphers



Queue Update Operations



53 / 121

Linked List Based

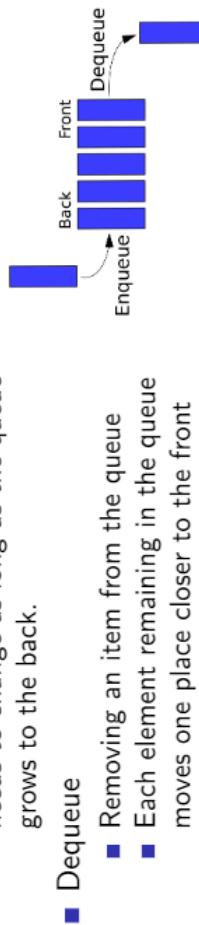
- 'head' is a pointer to the head of the list, 'tail' is a pointer to the end of the list
- Queue grows at the tail of the list

Pros of Linked List Based

- Internal storage is dynamically allocated → always right-sized

Cons of Linked List Based

- Enqueue/Dequeue have increased complexity



- Enqueue
 - Inserting an item into the queue
 - Imagine adding something to the end of the line. No other items' position needs to change as long as the queue grows to the back.
- Dequeue
 - Removing an item from the queue
 - Each element remaining in the queue moves one place closer to the front
 - To access an element behind, you must dequeue every element inserted before it
 - front() or peek() get the front item like dequeue() but without removing it from the queue

54 / 121

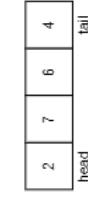
Queue Interface

Queue Implementations

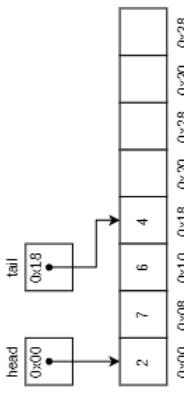


```
1  /* Not shown: new, delete, print methods */
2  Node *back(Queue *q);
3  void clear(Queue *q);
4  Node *dequeue(Queue *q);
5  Node *enqueue(Queue *q, Data d);
6  Node front(Queue *q);
7  size_t size(Queue *q);
8  void traverse(Queue *q, int (*func)(const void *));
```

Queue ADT



Array-Based Queue



59 / 121

57 / 121

Circular Buffer



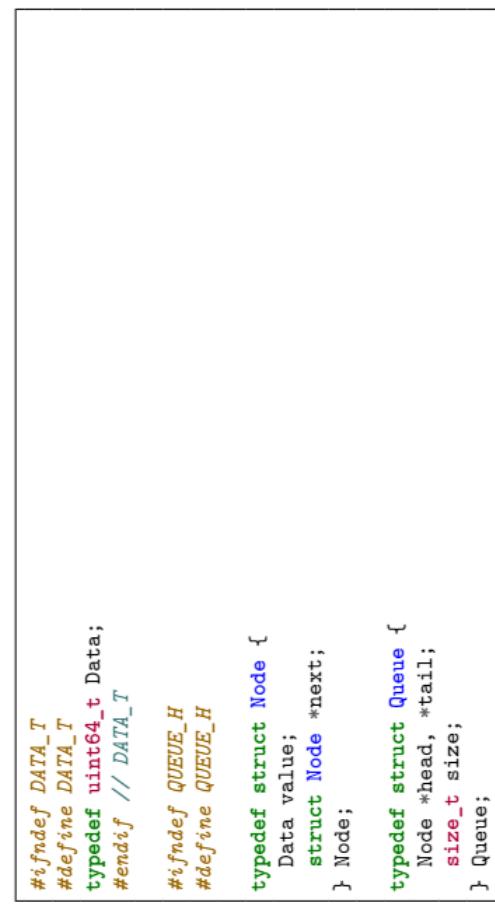
A specialized implementation of queues, using a single, fixed-size buffer with wraparound addressing as though the buffer was connected end-to-end. Like queues, circular buffers are also a sequential collection:

Sequential Collection

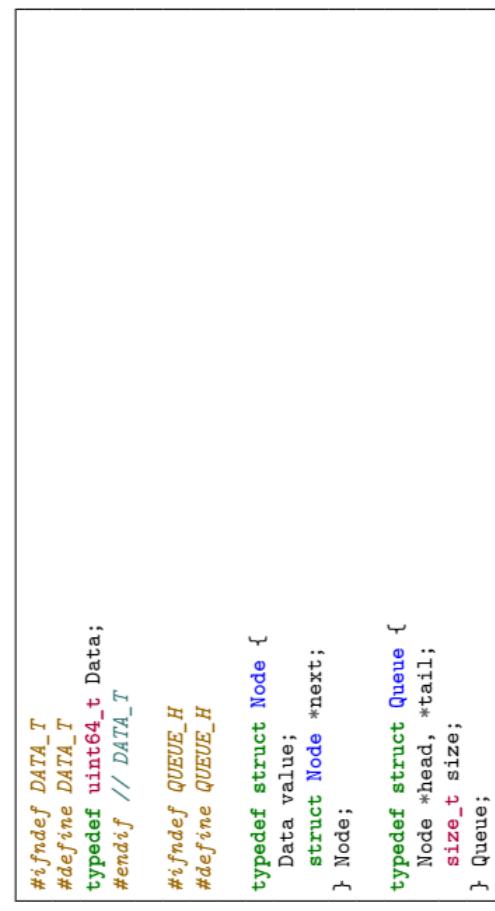
Contains items of the same type, forming a grouping through encapsulation

- Like queues, circular buffers provide restricted access control
- First In First Out (FIFO) access control, except they have the advantage of not needing any memory resizing or copying (all operations are constant time).

```
1  #ifndef DATA_T
2  #define DATA_T
3  typedef uint64_t Data;
4  #endif // DATA_T
5
6  #ifndef QUEUE_H
7  #define QUEUE_H
8  struct Node {
9      Data value;
10     struct Node *next;
11 } Node;
12
13
14  typedef struct Queue {
15     Node *head, *tail;
16     size_t size;
17 } Queue;
```



60 / 121



58 / 121

Examples of Circular Buffers

Circular Buffer Access Control



- Circular buffers are **First In, First Out (FIFO)**
- Items are removed in the order of their insertion.
- Think of circular buffers being like clocks, where the hands point to the item in storage to be read, or the location in storage to write the next item.

61 / 121

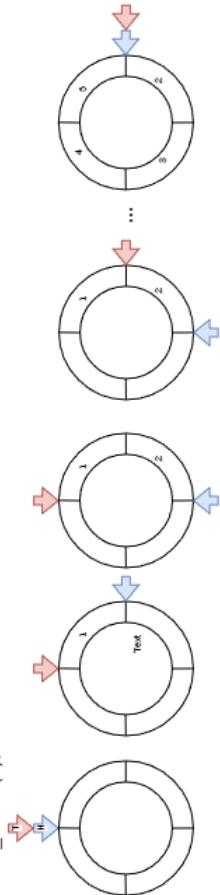
Circular Buffer Implementation



Examples of Circular Buffers



```
b = new_Buffer(4);    write(b,1);    write(b,2);    read(b);    write(b,5)
```



Write:

- 1 buffer[tail] ← value
- 2 tail ← tail + 1 mod capacity
- 3 return temp

Read:

- 1 temp ← buffer[tail]
- 2 tail ← tail - 1 mod capacity
- 3 return temp

What are some examples of circular buffers in computer science?

62 / 121



Circular Buffer Access Control



What are some examples of circular buffers in computer science?

- Keyboard input
- Producer-consumer pattern

63 / 121

62 / 121



- Deques can be implemented with different types of storage:
 - head-tail linked list (most common)
 - dynamic array (sometimes implemented as a circular array when maximum storage size is known)
- Deques can be implemented with different access restrictions:
 - Insertion and deletion at both ends (most common)
 - Input restricted: deletion from both ends, but insertion at one end only
 - Output restricted: insertion at both ends, but deletion at one end only

```

1 #include <stdint.h>
2 #include <stddef.h>
3
4 typedef uint64_t Data;
5 typedef struct Buffer {
6     size_t *head, *tail;
7     Data* storage;
8     size_t size;
9 }
10
11 // Not shown:
12 // new, copy, del methods for Buffer
13 void clear(Buffer *b);
14 void write(Buffer *b, Data d);
15 Data read(Buffer *b);

```

66 / 121

Examples of Deques



A **double-ended queue** or **deque** (pronounced *deck*) is a generalization of the queue plus the stack. Deques are still sequential collections:

- Sequential Collection** Contains items of the same type, forming a grouping
- Deques provide restricted access control through encapsulation
 - Insertion and removal can happen at either side of the storage, but not in the middle of the storage.

64 / 121

Double-Ended Queue

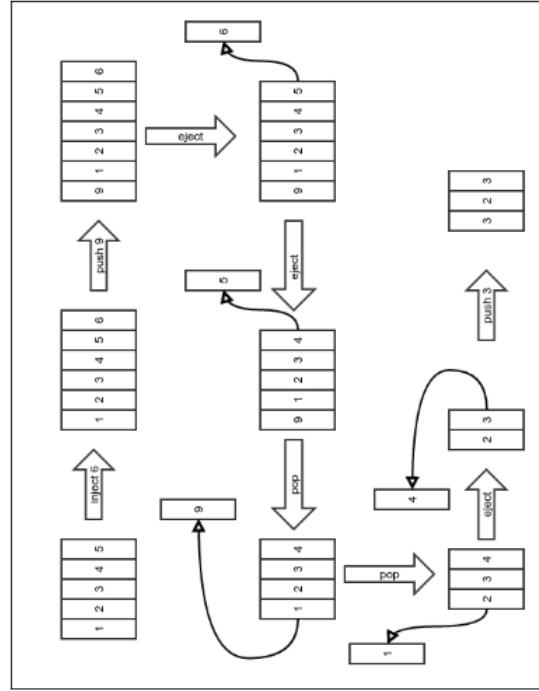


67 / 121

Deque Operation



Examples of Deques



69 / 121



Deque Operation



The basic operations of the deque are enqueue and dequeue at either end:

Enqueue:

- Insert element at back:
 - inject (textbooks)
 - push back (C-like languages)
 - ... or delete
- Insert element at front:
 - push (textbooks)
 - push front (C-like languages)
 - ... or prepend, unshift, cons

Dequeue:

- Remove element at back:
 - eject (textbooks)
 - pop back (C-like languages)
 - ... or delete
- Remove element at front:
 - pop (textbooks)
 - pop front (C-like languages)
 - ... or shift
- Deque implementations also sometimes offer a “peek” that returns the value without dequeuing it.

67 / 121

Trees

70 / 121

68 / 121

Properties of a Node



Can We Improve Our Existing ADTs?



key The value stored in a node.

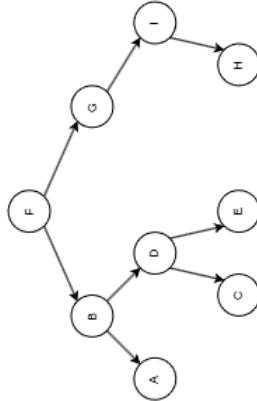
depth The distance of branches from a node to the root of its tree. For a root node, $depth = 0$.

level The set of nodes which have the same depth. The root node alone is $level = 0$

degree The number of children. For a leaf node, $degree = 0$.

Questions:

- What are the depths of A, C?
- What nodes are in level 3?
- What are the degrees of F, G, I and H?



73 / 121

Properties of a Node



Root Node The node with a depth of zero.

Leaf Node A node with no children.

Internal Node A node with one or more children.

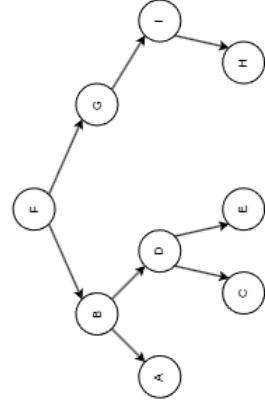
Sibling Node A node which shares a parent.

Ancestor Node All node(s) which are on the path between a node and the root.

Descendent Node All node(s) in a parent node's subtree.

True or False?

- Every tree has a single root node.
- Every non-empty tree has at least one leaf node.
- A root node can be a leaf node or an internal node, but not both.
- For any subtree, all nodes on a level have the same ancestors.



74 / 121

Definition of a Tree



Tree

A finite set of connected nodes such that one node is designated as the root, and all other nodes are recursively partitioned into sets called **subtrees**. Nodes are connected to one another by edges called **branches**.

Hierarchical Relationships Between nodes descend from **root node** to **leaf nodes**. There is exactly one path of **branches** from the root to each node. **Trees cannot be cyclic**.

Connected Each node except **root** must be connected to zero or more **one parent**, but may be connected to zero or more **children**. Trees cannot have more than one root.

72 / 121



Binary Search Tree (BST)

A tree with the additional properties:

- The tree's nodes have a *degree* ≤ 2
- Every parent node's left subtree has keys strictly less than the node's key, and the right subtree has keys strictly greater than the parent node's key.
- All node keys are unique.

BST Implementation

Array

- $index_{left} = 2(index_{parent}) + 1$
- $index_{right} = 2(index_{parent}) + 2$
- $index_{parent} = \lfloor \frac{index_{child}-1}{2} \rfloor$

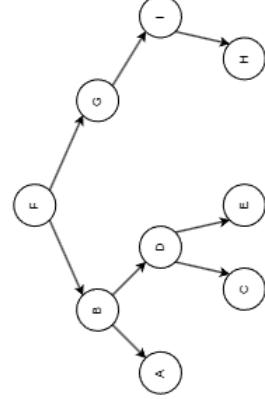
Pros

- $O(1)$ access, $O(n)$ storage

Cons

- Tree's max size must be known, or the array must be resized.
- How to represent empty nodes?

- **Size** The number of nodes in the tree.
- **Degree** The maximum degree of any node in the tree.
- **Height** The maximum depth of all nodes in the tree.
- **Width** The number of nodes in a level.
- **Breadth** The number of leaves in the tree.



75 / 121

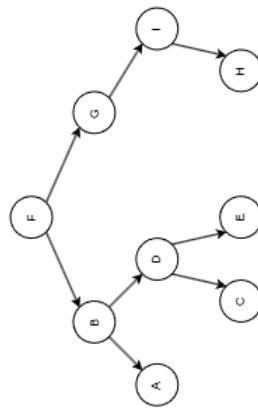


Properties of a Tree

Full Every node has either a zero or the maximum degree

Complete All levels are full, except for possibly the last level, and the nodes in that level are as far left as possible.

Balanced A tree whose left and right subtrees' heights differ by not more than one level.



- Is this tree full?
- Is this tree complete?
- Is this tree balanced?

78 / 121

76 / 121

$$f(n) = \begin{cases} \text{atomic object} & \text{base case} \\ f(n') & \text{otherwise} \end{cases}$$

- Recursion is process of a function calling itself to iterate →
- Use the stack as loop
- **How does it work?**

- 1 Determine *base case* that ends recursion and produces an *atomic object*
- 2 Determine recursive step that determines how objects can be combined to produce another object closer to the atomic object
- 3 Every recursive step must make progress towards base case:

- **cannot use circular definitions or we get infinite recursion**
- **Why use recursion?** Simplifies producing code for repetitive operations
- **Why not use recursion?** Uses more memory, unwinding can be slow – in some cases may re-calculate known values.

81 / 121

79 / 121

Recursion and Trees



Property, access and update functions use recursion. Examples:

- $\text{size(tree)} = \text{size(left subtree)} + \text{size(right subtree)} + 1 \rightarrow$
Base case: empty subtree has size 0
- $\text{height(tree)} = \max(\text{height(left subtree)}, \text{height(right subtree)}) + 1 \rightarrow$ Base case: empty subtree has size 0
- $\text{min(tree)} = \min(\text{left subtree}) \rightarrow$ Base case: min value of subtree of a leaf node is leaf node's value.
- Insert, deletion, search, traversal all easiest to implement using recursive operations on subtrees

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15		
50	17	72	12	23	54	76	9	14	19		67						

80 / 121

81 / 121

Search



Insertion



- Almost identical to insert, except we return reference to found node.
- Use recursion to check if value is greater, lesser, or equal to the current node
 - If lesser: go to left subtree, check again
 - If greater: go to right subtree, check again
 - If equal: return reference
 - If null: value not found, return nullptr

```
recursiveFind(node, data)
    if(data == node.data)
        return node.data
    else if (data < node.data)
        if(node.left == NULL)
            return NULL;
        else
            return read(node.left, data)
    else if (data > node.data)
        if(node.right == NULL)
            return NULL;
        else
            return read(node.right, data)
```

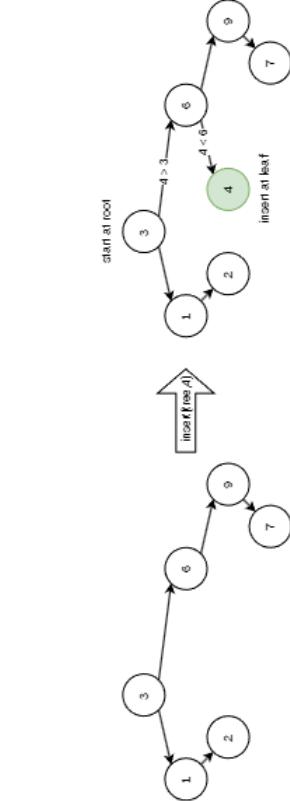
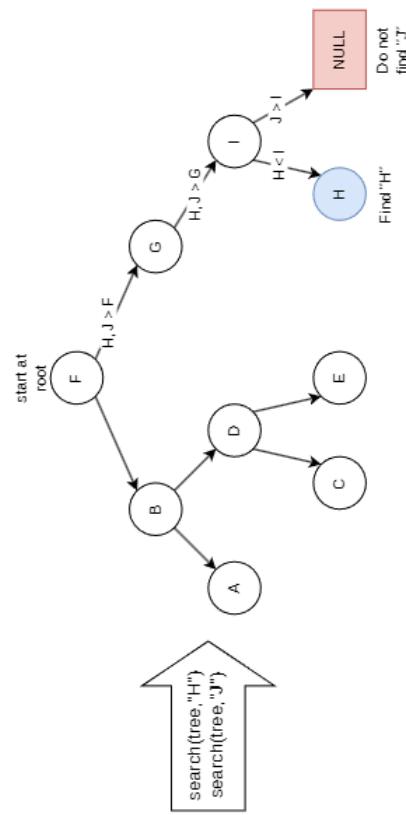
85 / 121

- Always start recursion at root, insert at leaf node.
- If tree is empty, inserted value becomes root node.
- Use recursion to check if value is greater, lesser, or equal to the current node
 - If lesser: go to left subtree, check again
 - If greater: go to right subtree, check again
 - If equal: value is already in tree, no duplicates allowed

```
recursiveInsert(node, data)
    if (data < node.data)
        if(node.left == NULL)
            addLeaf(node, data)
        else
            insert(node.left, data)
    else if (data > node.data)
        if(node.right == NULL)
            addLeaf(node, data);
        else
            insert(node.right, data)
```

83 / 121

Insertion



```

1 removeLeaf(node)
2 // disconnect parents or empty tree
3 if (node is right child)
4     node->parent->right = NULL
5 else if (node is left child)
6     node->parent->left = NULL
7 else (node is root)
8     tree->root = NULL
9
10 // delete the node
11 node->left,right,parent = NULL
12 return node

```

Basic deletion process:

- Easiest case:
 - disconnect from parent, return pointer to the deleted leaf.
 - What if the leaf is also the tree's root?
 - Set tree's root to NULL signifying an empty tree.
 - If the leaf is not the root, then it must have a parent. We can then follow the parent pointers up the tree to find the root.

89 / 121

```

shortCircuit(node)
if (node is right child)
    if (node has right child)
        node->parent->right = node->right
    else (node has left child)
        node->parent->right = node->left
    else if (node is left child)
        if (node has right child)
            node->parent->left = node->right
        else (node has left child)
            node->parent->left = node->left
    // delete the node
    node->left,right,parent = NULL
    return node

```

```

node = recursiveFind(tree->root, data)
if (node not in BST)
    return NULL
else if (node has 0 subtrees)
    return deleteLeaf(node)
else if (node has 1 subtree)
    return shortCircuit(node)
else if (node has 2 subtrees)
    return promotion(node)

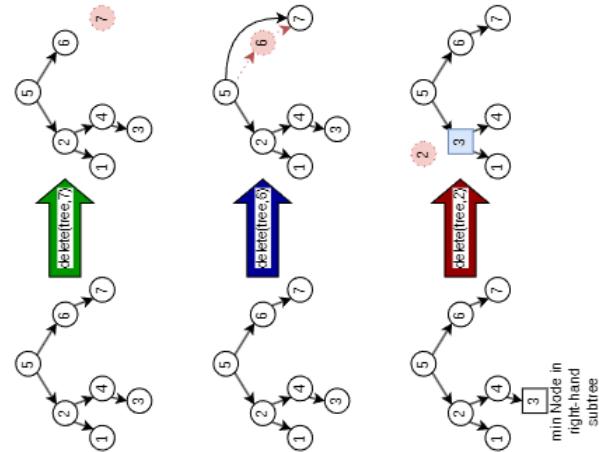
```

When deleting a parent node in BST, ensure that the BST ordering property is preserved!

Basic deletion process:

- Find the node to be removed by its value
 - Unlink the node from the tree
 - Re-link the removed node's parent/children if necessary
 - Return a reference to the removed node and/or free the node as required
- Every deletion is one of three cases:**
- Delete a leaf node → easiest case
 - Delete a interior node with order 1 → can't just delete node because then the tree becomes un-connected.
 - Delete an interior node with order 2 → must choose which child becomes the new parent

87 / 121



98 / 121

Traversal

Deletion promotion()



Traversal

A specialized form of a graph walk that begins at a tree's root node, and visits each node exactly once. Traversals are categorized by the order in which the nodes are visited.

Two types of traversal:

Depth-first traversal The search tree is deepened as far as possible before visiting a sibling node. Deferred visits are stored in a LIFO data structure such as a stack.

Breadth-first traversal The search tree is broadened as far as possible before visiting a child level. Deferred visits are stored in a FIFO data structure such as a queue.

Depth First Traversal

Deletion promotion()

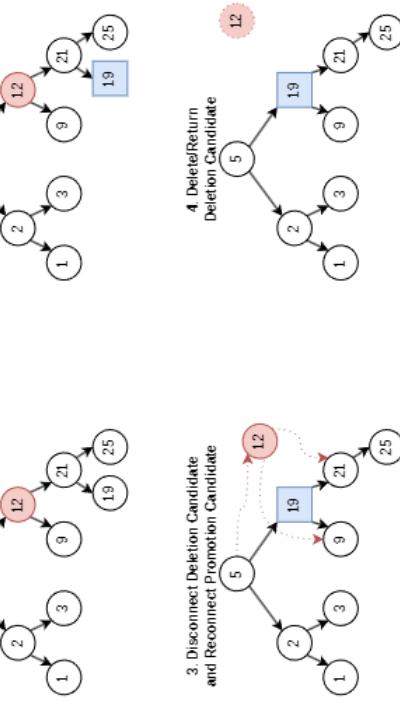


- We must promote a node to a higher space in the tree to preserve BST connectedness property!
- There are two possible "successor" candidates, it doesn't matter which is picked, both preserve BST ordering property:
 - max node in the left subtree → Traverse left once, then right as far as possible
 - min node in the right subtree → Traverse right once, then left as far as possible
- Successor candidate may be more than one level below the deletion candidate!
- What if subtree min or max is not a leaf node? Use short circuit to remove it and promote its child in place.
- What if deletion candidate is also the tree's root node?
 - Successor candidate must also become tree's new root node.

93 / 121

Basic Algorithm:

- 1 Start at root node
- 2 Recursively perform the following operations:
 - 1 If the current node is empty, then return.
 - 2 Otherwise, perform the following operations in a certain order:
 - L: Recursively traverse the current node's left subtree
 - N: Visit the current node
 - R: Recursively traverse the current node's right subtree
- Visit order is described as **Pre-order** (NLR/NRL), **In-order** (LNR/RNL), or **Post-order** (LRN/RLN)



91 / 121



Deletion promotion()



- Note that we do not simply swap the values, we actually disconnect/existing nodes.
- In some cases, value swap would work. But in our project test suite, it will cause a "possible" memory leak according to valgrind's analysis if you aren't meticulous.

94 / 121

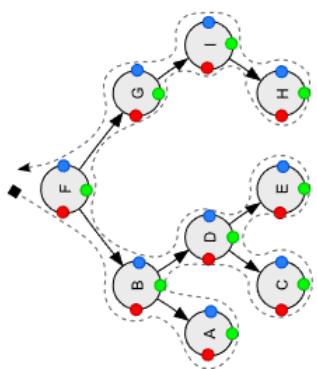
95 / 121

Depth-first Traversal

Post-order (LRN)

Depth-first Traversal

Pre-order (NLR)



```

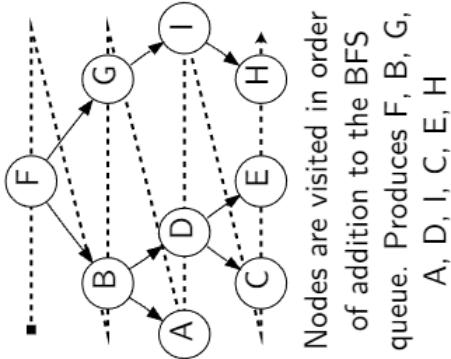
1 RecursivePostorder(node):
2   RecursivePostorder(node->left)
3   RecursivePostorder(node->right)
4   Visit(node)

```

Following blue dots
produces A, C, E, D, B, H,
I, G, F → generates
reverse topological sort
and used to delete the tree

Breadth-first Traversal

Level-order



```

1 RecursiveBFS(tree):
2   Queue q = empty
3   Enqueue(q, tree->root)
4   while q is not empty:
5     node = Dequeue(q)
6     Visit(node)
7     for child in node->children:
8       Enqueue(q, child)

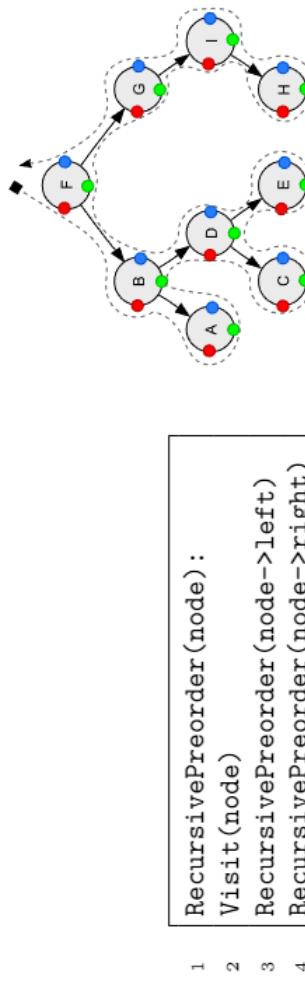
```

Nodes are visited in order
of addition to the BFS
queue. Produces F, B, G,
A, D, I, C, E, H



Depth-first Traversal

In-order (LNR)

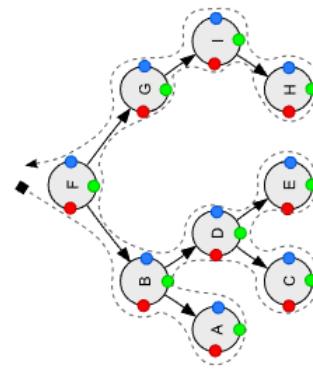


```

1 RecursiveInorder(node):
2   RecursiveInorder(node->left)
3   Visit(node)
4   RecursiveInorder(node->right)

```

Following red dots
produces F, B, A, D, C, E,
G, I, H → *topological sort*
used for copying the tree



```

1 RecursiveInorder(node):
2   RecursiveInorder(node->left)
3   Visit(node)
4   RecursiveInorder(node->right)

```

Following green dots
produces A, B, C, D, E, F,
G, H, I → *lexicographical sort*
used to give nodes in
order of tree's underlying
comparison function and
to balance the tree



Array Based Heaps



Recall that we can map elements of a tree into a zero-indexed array:

```
function LEFT(i)           ▷ The left child of the parent at i
    return 2i + 1
end function

function RIGHT(i)          ▷ The right child of the parent at i
    return 2i + 2
end function

function PARENT(i)          ▷ The parent of the child at index i
    return ⌊ $\frac{i-1}{2}$ ⌋
end function
```

- The priority member is always at $A[0]$
- There will never be an unfilled node at a smaller index than a filled node's index because the heap is *complete*.

101 / 121

Array Based Heaps



Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	16	14	10	8	7	9	3	2	4	1					

Heap Definition



Heap

A partially-ordered, complete tree that satisfies the *heap property*.

Partially Ordered Ancestors and descendants have an ordered relationship according to the *heap property*.

- The heap is not sorted. There is no ordered traversal of the nodes.
- There is no ordering between two sibling or cousin nodes like in a BST.

Complete All levels are full, except for possibly the last level, and the nodes in that level are as far left as possible.

Heap Property For any given node C and its' parent node P, then

- In a *max heap*, $\text{key}(P) \geq \text{key}(C)$.
- In a *min heap*, the $\text{key}(P) \leq \text{key}(C)$.

102 / 121

100 / 121

Maintaining the Heap Property

Maintaining the Heap Property



```
1: function HEAPIFY( $A, i$ )
```

```
2:    $I \leftarrow \text{LEFT}(i)$ 
```

```
3:    $r \leftarrow \text{RIGHT}(i)$ 
```

```
4:   if  $I < \text{size}[A]$  and  $A[I] > A[i]$  then
```

```
5:      $\text{largest} \leftarrow I$   $\triangleright$  largest is index of max key in parent-children subheap
```

```
6:   else
```

```
7:      $\text{largest} \leftarrow i$ 
```

```
8:   end if
```

```
9:   if  $r < \text{size}[A]$  and  $A[r] > A[i]$  then
```

```
10:     $\text{largest} \leftarrow r$ 
```

```
11:   end if
```

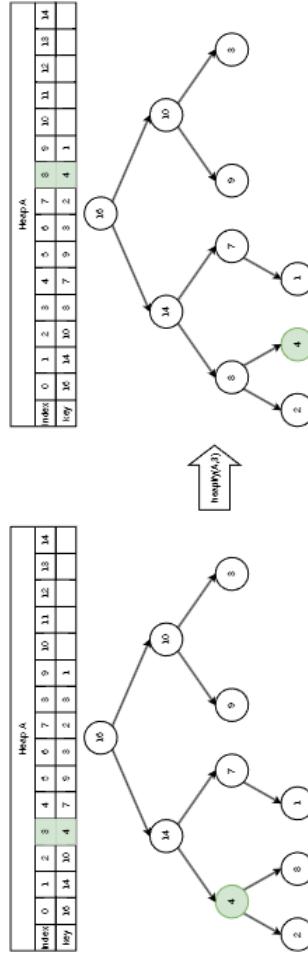
```
12:   if  $\text{largest} \neq i$  then
```

```
13:     exchange  $A[i] \leftrightarrow A[\text{largest}]$ 
```

```
14:     HEAPIFY( $A, \text{largest}$ )
```

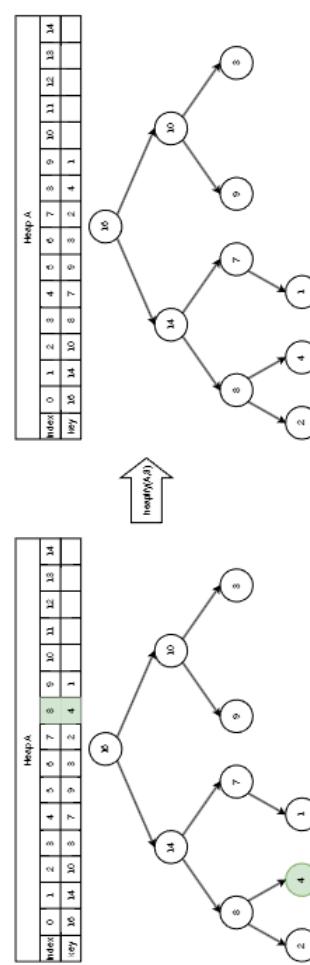
```
15:   end if
```

```
16: end function
```



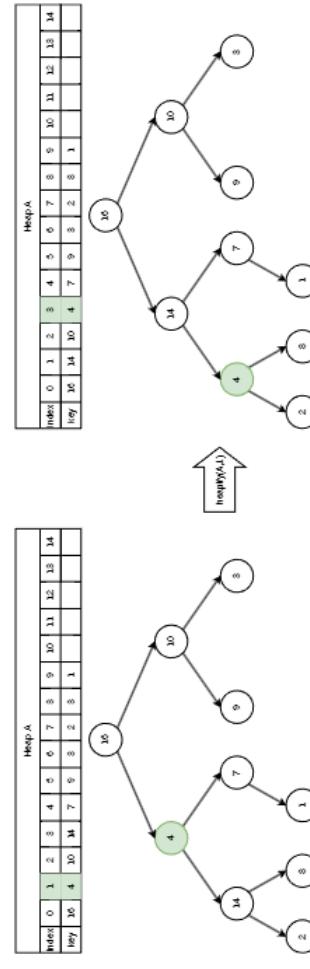
104 / 121

Maintaining the Heap Property



103 / 121

Maintaining the Heap Property



104 / 121

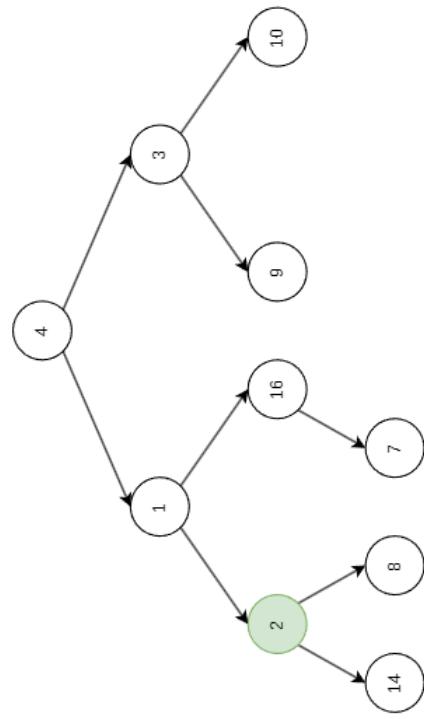
104 / 121

Building a Heap

Building a Heap



Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	4	1	3	2	16	9	10	14	8	7					



106 / 121

Heap A

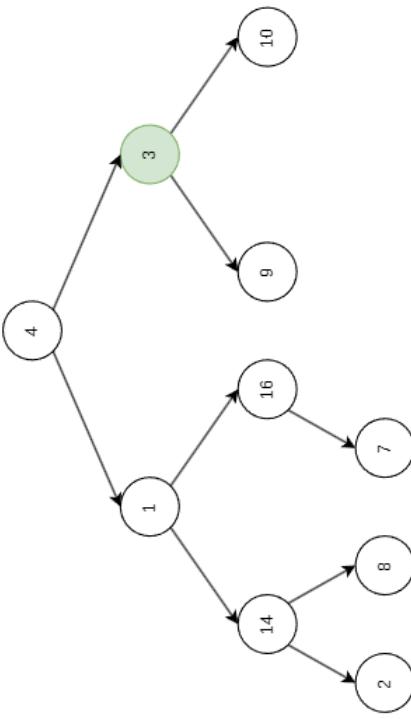
```
1: function BUILD-HEAP(A)           ▷ Convert array to heap object
2:   heapsize[A] ← length[A]        ▷ index of last interior node
3:   n ← ⌊length[A] - 1⌋
4:   for i ← n downto 0 do
5:     HEAPIFY(A, i)               ▷ "Sift Down"
6:   end for
7: end function
```

105 / 121

Building a Heap



Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	4	1	3	14	16	9	10	2	8	7					



106 / 121

Heap A

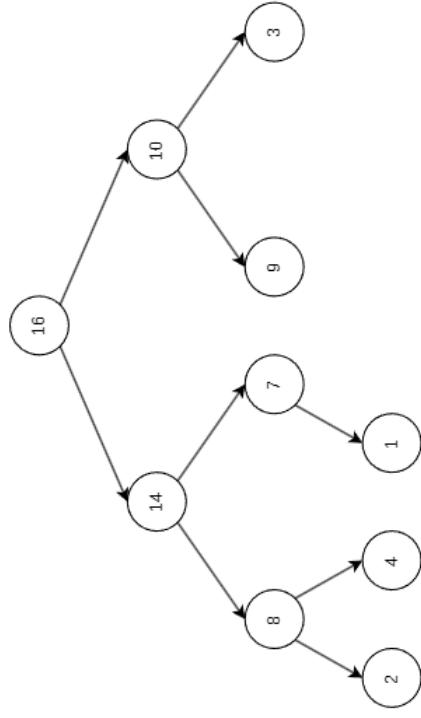
Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	4	1	3	14	16	9	10	2	8	7					

Building a Heap

Building a Heap

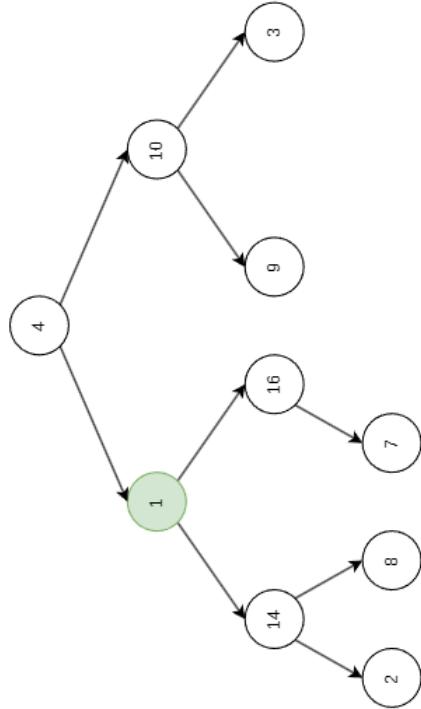


Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	16	14	10	8	7	9	3	2	4	1					



106 / 121

Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	4	1	10	14	16	9	3	2	8	7					



106 / 121

Sorting an Array using Heaps

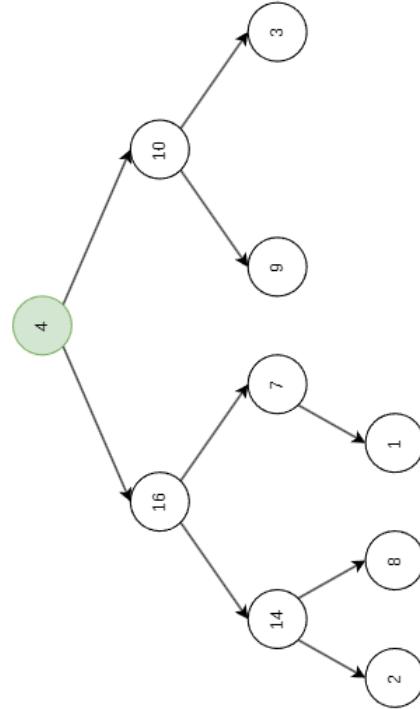


Building a Heap

Building a Heap



Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	4	16	10	14	7	9	3	2	8	1					



```
1: function HEAPSORT(A)
2:   length ← heapsize[A]
3:   for i ← length downto 1 do
4:     exchange A[0] ↔ A[i]
5:     heapsize[A] ← heapsize[A] - 1
6:     HEAPIFY(A, 0)
7:   end for
8:   heapsize[A] ← length
9:   return A ▷ array in reverse priority order, no longer a heap
10: end function
```

107 / 121

106 / 121

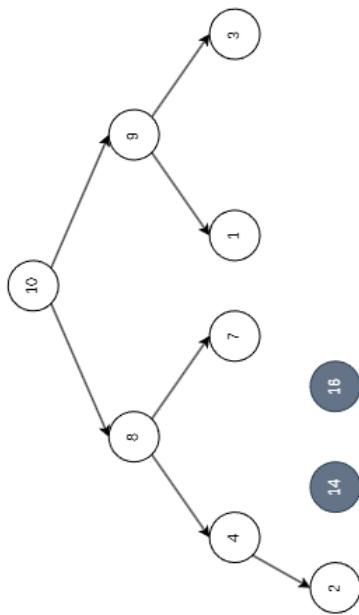
Sorting an Array using Heaps



Sorting an Array using Heaps



Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	10	8	9	4	7	1	3	2	14	16					



108 / 121

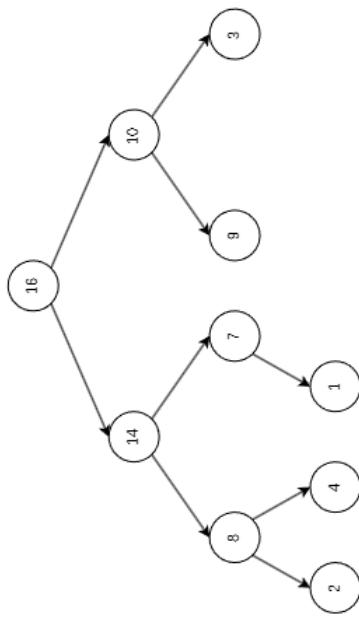
Sorting an Array using Heaps



Sorting an Array using Heaps

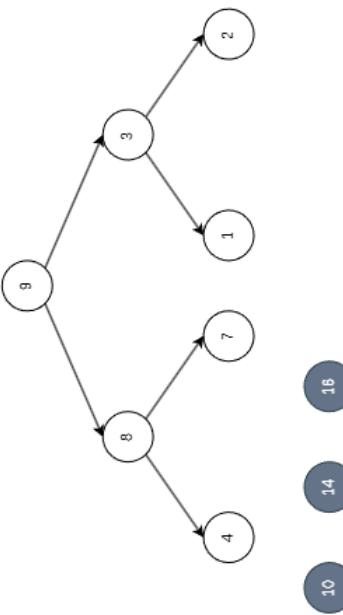


Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	16	14	10	8	7	9	3	2	4	1					



108 / 121

Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	14	8	10	4	7	9	3	2	1	16					



108 / 121

Heap A															
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
key	14	8	10	4	7	9	3	2	1	16					



Sorting an Array using Heaps

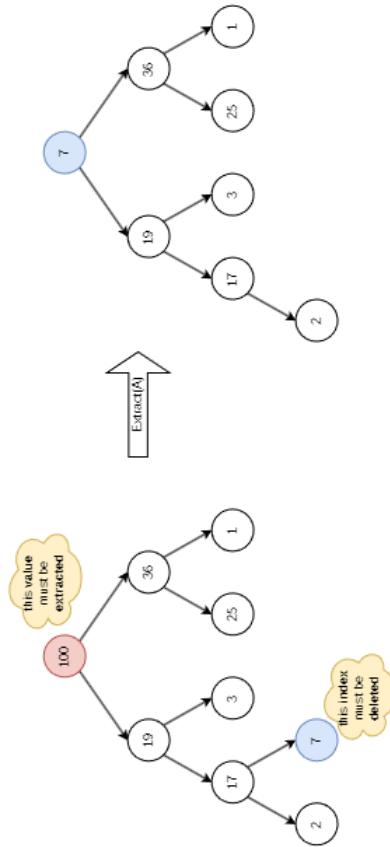


Sorting an Array using Heaps

108 / 121

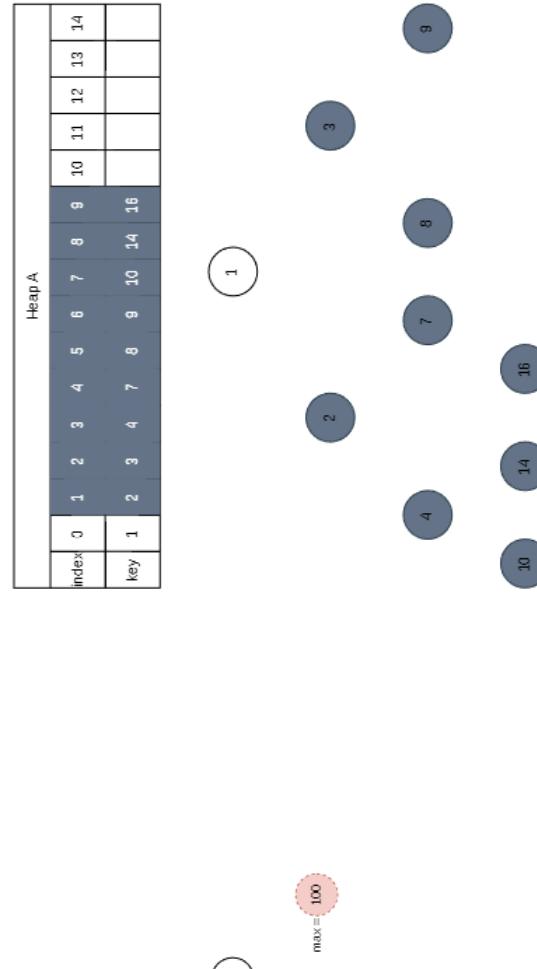
Getting the “Best” Element

Sorting an Array using Heaps



110 / 121

Getting the “Best” Element



108 / 121

Getting the “Best” Element

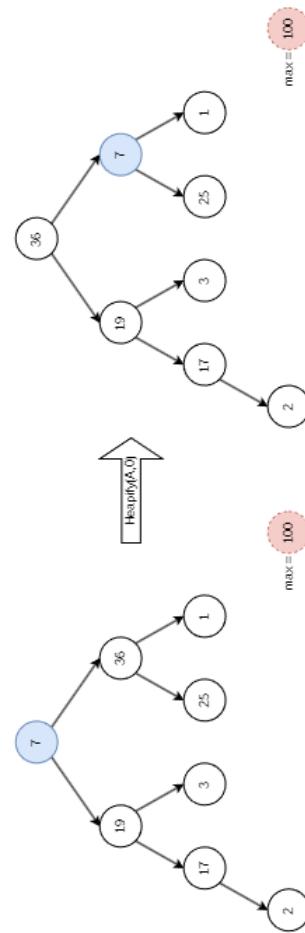


```
1: function MAXIMUM(A)
2:   return A[0]
3: end function

1: function EXTRACT(A)
2:   if heapsize[A] < 1 then
3:     error "heap underflow"
4:   end if
5:   max ← A[0]
6:   A[0] ← A[heapsize[A] - 1]
7:   heapsize[A] ← heapsize[A] - 1
8:   HEAPIFY(A, 0)
9:   return max
10: end function
```

▷ “Sift Down”

110 / 121



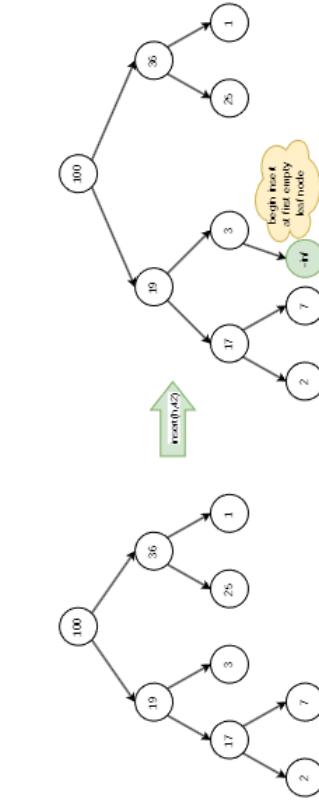
109 / 121

Updating and Inserting Keys

Getting the “Best” Element



Getting the “Best” Element

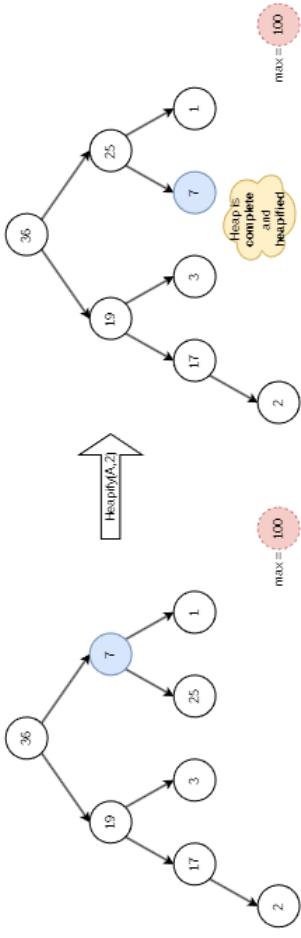


112 / 121

Updating and Inserting Keys



Updating and Inserting Keys



110 / 121

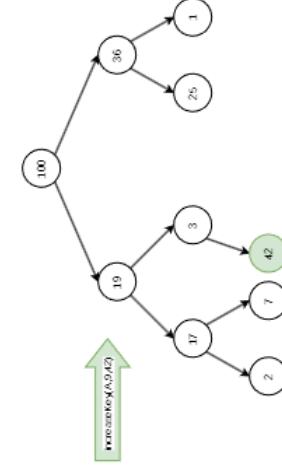
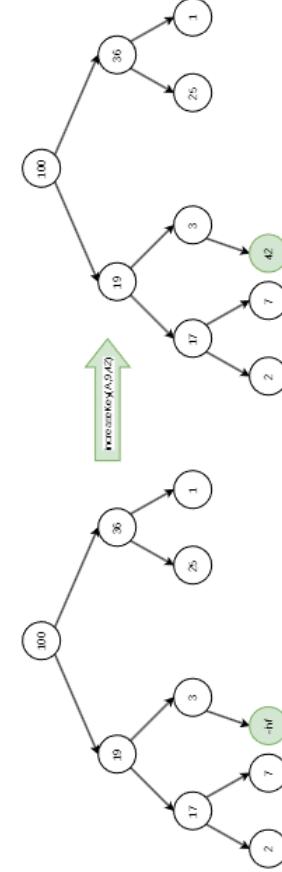


Updating and Inserting Keys

```
1: function INCREASEKEY( $A, i, key$ )
2:   if  $key < A[i]$  then
3:     error "new key is smaller than old key"
4:   end if
5:    $A[i] \leftarrow key$ 
6:   while  $i > 0$  and  $A[PARENT(i)] < A[i]$  do
7:     exchange  $A[i] \leftrightarrow A[PARENT(i)]$           ▷ "Sift Up"
8:      $i \leftarrow PARENT(i)$ 
9:   end while
10: end function
```

```
1: function INSERT( $A, key$ )
2:    $heapsize[A] \leftarrow heapsize[A] + 1$ 
3:    $A[heapsize[A] - 1] \leftarrow -\infty$ 
4:   INCREASEKEY( $A, heapsize[A] - 1, key$ )
5: end function
```

112 / 121



111 / 121

Priority Queue



Updating and Inserting Keys



A specialized implementation of queues, priority queues are queues where insertion location is based on a *priority value*. Priority queues are a sequential collection:

Sequential Priority of elements causes ordering

Collection Contains items of the same type, forming a grouping

■ Priority queues are an ADT providing restricted access control through encapsulation

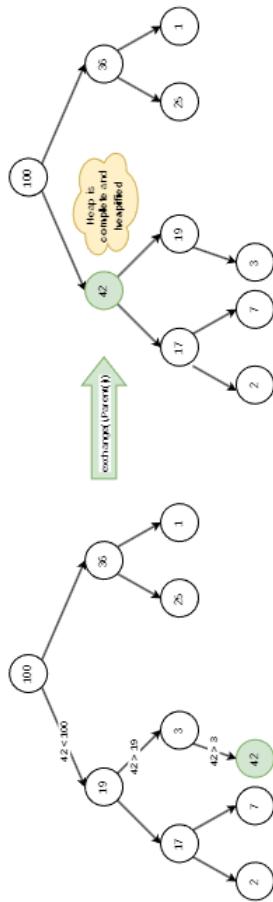
■ Like queues, only one element can be removed (queue: first in, priority queue: highest priority).

■ Unlike queues, insertion happens anywhere in the priority queue, not just the tail.

■ Don't confuse an item's *priority* with its *value*
■ Consider a priority queue in a scheduler:

→ value is the process' PID
→ priority is the preference to schedule that process (least recently run, lowest "niceness", etc.)

114 / 121



112 / 121

Priority Queue Access Control



■ Queues are **First In, First Out (FIFO)** ...

■ But priority queues are
"Best" Priority Served First

■ Priority queues implement operations congruent to sorting algorithms

■ Unlike queues, priority queues are often not built with vectors, but heaps or binary search trees!

■ Think of priority queues being like airport queues where pilots and flight crew



115 / 121

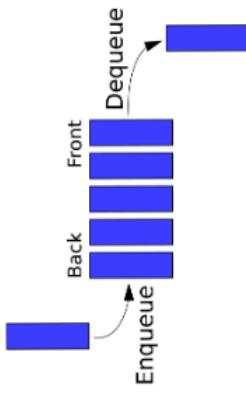
Priority Queues



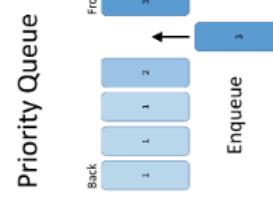
113 / 121

Queue vs Priority Queue

Examples of Priority Queues



What are some examples of priority queues in computer science?



Design Considerations



Choosing an underlying ADT implementation can have a dramatic effect on the performance of a priority queue.

- Array
 - Sorted array makes insertion slow, but removal fast
 - Unsorted array makes insertion fast, but removal slow
- Linked List
 - Sorted list makes insertion slow, but removal fast
 - Unsorted list makes insertion fast, but removal slow
- Binary Search Tree
 - Insertion and removal are both fast as long as the tree is balanced
 - Duplicate keys not possible – useful for set-like queues which enforce uniqueness.

Examples of Priority Queues



What are some examples of priority queues in computer science?

- Priority scheduling (as implemented in early Linux and XV6)
- Bandwidth management in network routing
- Huffman coding for data compression
- Efficient path finding algorithms (A* algorithm)
- Efficient spanning algorithms (Prim's algorithm)



Thought questions:

- If the priority item is always at the root, why isn't removing it constant time?
- Why can't we search for a particular item quickly if the heap property is conserved?
- Why are duplicates a problem for heap-based priority queues?
- What changes would we have to make to change from a max-priority queue to a min-priority queue?

Thought questions:

- If the Priority Queue's storage list is singly or doubly linked?
- What inputs would cause a Binary Search Tree to be unbalanced?
- If the tree isn't balanced, why is performance is basically equivalent to a sorted linked list?

Heap-Based Priority Queue



- Insertion and removal of the priority item are both logarithmic time.
- Inspection (without removal) of the priority item is constant time.
- Search for a particular key is linear time and generally not supported.
- Duplicate keys permitted, but no way to resolve "ties" in priority without refactoring our removal method.